

02616

# Large-scale Modelling

recap week 4

`mpi4py`

# MPI

- How Recv/Send are used
  - Recv is a *maximum* buffer size!
- The difference between send/Send (prefer Send)
  - Send/Recv is for pre-allocated numpy buffers
  - send/recv is for Python objects (slow and lots of overhead)
- How to use the buffer argument to control amount of data sent/received
  - `comm.Send(buf, ...)` # auto dtype + count
  - `comm.Send([buf, 1], ...)` # auto dtype
  - `comm.Send([buf, MPI.INT], ...)` # auto count
  - `comm.Send([buf, 1, MPI.INT], ...)` # explicit
- How to use non-blocking communications

```
req = comm.Isend(buf, ...)
# do something else than writing to buf
req.Wait() # buf is now sent!
```

  - Remember to *always* post a Wait/Test for all requests

# MPI

- How to use Status for getting data

```
comm.Recv(..., status=status)
count = status.Get_count(MPI.DOUBLE)
tag = status.tag
source = status.source
```

- How to use *collectives*

```
#rank=0: [0, 1] → [0, 1]
#rank=1:           → [0, 1]
comm.Bcast(buf, root=0)
#rank=0: [0, 1] → [0]
#rank=1:           → [1]
comm.Scatter(buf_send, buf_recv, root=0)
#rank=0: [0] → [0, 1]
#rank=1: [1] →
comm.Gather(buf_send, buf_recv, root=0)
#rank=0: [0, 4] → [0, 3]
#rank=1: [1, 3] →
comm.Reduce(buf_send, buf_recv, op=MPI.MIN,
            root=0)
```